


About the Project

This project explores coordinate-based neural networks for image reconstruction. A fully connected multilayer perceptron is implemented from scratch using NumPy, where the input is a 2D pixel coordinate (x, y) and the output is the corresponding RGB color value (r, g, b) .



To investigate how the network captures spatial and frequency details, I compared several input feature mapping strategies $\gamma(v)$:

- No mapping: $\gamma(v) = v$.
- Basic mapping: $\gamma(v) = [\cos(2\pi v), \sin(2\pi v)]^T$.
- Gaussian Fourier feature mapping: $\gamma(v) = [\cos(2\pi Bv), \sin(2\pi Bv)]^T$, where each entry in $B \in \mathbb{R}^{m \times d}$ is sampled from $N(0, \sigma^2)$.

Implementation Highlights:

- Mappings are implemented in the helper functions `get_B_dict` and `input_mapping`.
- The basic mapping can be considered a case where $B \in \mathbb{R}^{2 \times 2}$ is the identity matrix.
- For this project, d is 2 because the input coordinates in two dimensions.

Setup

```
from google.colab import drive
drive.mount("/content/drive")

import os
datadir = "path/"
if not os.path.exists(datadir):
    !ln -s "path/" $datadir
os.chdir(datadir)
!pwd
```

Imports

```
import matplotlib.pyplot as plt
from tqdm.notebook import tqdm
import os, imageio
import cv2
import numpy as np
```

```

from base64 import b64encode
from IPython.display import HTML

from models.neural_net import NeuralNetwork

# makes sure your NeuralNetwork updates as you make changes to the .py file
%load_ext autoreload
%autoreload 2

# sets default size of plots
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0)

The autoreload extension is already loaded. To reload it, use:
%reload_ext autoreload

```

Helper Functions

Image Data and Feature Mappings

```

# Data loader - already done for you
def get_image(size=512, \

image_url='https://bmild.github.io/fourfeat/img/lion_orig.png'):

    # Download image, take a square crop from the center
    img = imageio.imread(image_url)[..., :3] / 255.
    c = [img.shape[0]//2, img.shape[1]//2]
    r = 256
    img = img[c[0]-r:c[0]+r, c[1]-r:c[1]+r]

    if size != 512:
        img = cv2.resize(img, (size, size))

    plt.imshow(img)
    plt.show()

    # Create input pixel coordinates in the unit square
    coords = np.linspace(0, 1, img.shape[0], endpoint=False)
    x_test = np.stack(np.meshgrid(coords, coords), -1)
    test_data = [x_test, img]
    train_data = [x_test[:, ::2, ::2], img[:, ::2, ::2]]

    return train_data, test_data

# Create the mappings dictionary of matrix B - you will implement this
def get_B_dict(size, scale = 1.0):
    mapping_size = size // 2

```

```

# mapping_size = size (for extra credit)
B_dict = {}
B_dict['none'] = None

# Basic mapping
basic_mapping = np.eye(2)
B_dict['basic'] = basic_mapping

# Gaussian mapping
gauss_mapping = np.random.normal(scale=scale, size =
(mapping_size, 2))
B_dict['gauss_1.0'] = gauss_mapping

return B_dict

# Given tensor x of input coordinates, map it using B
def input_mapping(x, B):
    if B is None:
        return x
    else:
        proj = 2 * np.pi * (x @ B.T)
        cos_term = np.cos(proj)
        sin_term = np.sin(proj)
        mapped_features = np.concatenate([cos_term, sin_term], axis=-
1)
        return mapped_features

```

MSE Loss and PSNR Error

```

# Loss function and evaluation metric
def mse(y, p):
    return np.mean((y - p) ** 2)

def psnr(y, p):
    return -10 * np.log10(2.*mse(y, p))

# Display image
size = 32
train_data, test_data = get_image(size)

/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting
with ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
    img = imageio.imread(image_url)[..., :3] / 255.

```



```
# Set the hyperparameters
num_layers = 4

hidden_size = 256
epochs = 1000
learning_rate = 0.1
output_size = 3
B_dict = get_B_dict(size)

print('B_dict items:')
for k,v in B_dict.items():
    print('\t',k,np.array(v).shape)
```

```

B_dict items:
  none ()
  basic (2, 2)
  gauss_1.0 (16, 2)

# Apply the input feature mapping to the train and test data
def get_input_features(B_dict, mapping):
    # mapping is the key to the B_dict, which has the value of B
    # B is then used with the function `input_mapping` to map x
    y_train = train_data[1].reshape(-1, output_size)
    y_test = test_data[1].reshape(-1, output_size)
    X_train = input_mapping(train_data[0].reshape(-1, 2),
B_dict[mapping])
    X_test = input_mapping(test_data[0].reshape(-1, 2), B_dict[mapping])
    return X_train, y_train, X_test, y_test

```

Plotting and video helper functions

```

def plot_training_curves(train_loss, train_psnr, test_psnr):
    # plot the training loss
    plt.subplot(2, 1, 1)
    plt.plot(train_loss)
    plt.title('MSE history')
    plt.xlabel('Iteration')
    plt.ylabel('MSE Loss')

    # plot the training and testing psnr
    plt.subplot(2, 1, 2)
    plt.plot(train_psnr, label='train')
    plt.plot(test_psnr, label='test')
    plt.title('PSNR history')
    plt.xlabel('Iteration')
    plt.ylabel('PSNR')
    plt.legend()

    plt.tight_layout()
    plt.show()

def plot_reconstruction(p, y):
    p_im = p.reshape(size,size,3)
    y_im = y.reshape(size,size,3)

    plt.figure(figsize=(12,6))

    # plot the reconstruction of the image
    plt.subplot(1,2,1), plt.imshow(p_im), plt.title("reconstruction")

    # plot the ground truth image
    plt.subplot(1,2,2), plt.imshow(y_im), plt.title("ground truth")

```

```

print("Final Test MSE", mse(y, p))
print("Final Test psnr", psnr(y, p))

def plot_reconstruction_progress(predicted_images, y, N=8):
    total = len(predicted_images)
    step = total // N
    plt.figure(figsize=(24, 4))

    # plot the progress of reconstructions
    for i, j in enumerate(range(0, total, step)):
        plt.subplot(1, N+1, i+1)
        plt.imshow(predicted_images[j].reshape(size, size, 3))
        plt.axis("off")
        plt.title(f"iter {j}")

    # plot ground truth image
    plt.subplot(1, N+1, N+1)
    plt.imshow(y.reshape(size, size, 3))
    plt.title('GT')
    plt.axis("off")
    plt.show()

def plot_feature_mapping_comparison(outputs, gt):
    # plot reconstruction images for each mapping
    plt.figure(figsize=(24, 4))
    N = len(outputs)
    for i, k in enumerate(outputs):
        plt.subplot(1, N+1, i+1)
        plt.imshow(outputs[k]['pred_imgs'][-1].reshape(size, size, -1))
        plt.title(k)
    plt.subplot(1, N+1, N+1)
    plt.imshow(gt)
    plt.title('GT')
    plt.show()

    # plot train/test error curves for each mapping
    iters = len(outputs[k]['train_psnrs'])
    plt.figure(figsize=(16, 6))
    plt.subplot(121)
    for i, k in enumerate(outputs):
        plt.plot(range(iters), outputs[k]['train_psnrs'], label=k)
    plt.title('Train error')
    plt.ylabel('PSNR')
    plt.xlabel('Training iter')
    plt.legend()
    plt.subplot(122)
    for i, k in enumerate(outputs):
        plt.plot(range(iters), outputs[k]['test_psnrs'], label=k)
    plt.title('Test error')
    plt.ylabel('PSNR')

```

```
plt.xlabel('Training iter')
plt.legend()
plt.show()
```

Experiment Runner

```
def NN_experiment(X_train, y_train, X_test, y_test, input_size,
num_layers,\
                hidden_size, output_size, epochs,\
                learning_rate, loss_function, opt='SGD', b1=0.9,
b2=0.999, eps=1e-8,
                batch_size = 32):

    # Initialize a new neural network model
    hidden_sizes = [hidden_size] * (num_layers - 1)
    current_lr = learning_rate
    net = NeuralNetwork(input_size, hidden_sizes, output_size,
num_layers, opt, loss_function)

    # Variables to store performance for each epoch
    train_loss = np.zeros(epochs)
    train_psnr = np.zeros(epochs)
    test_psnr = np.zeros(epochs)
    predicted_images = np.zeros((epochs, y_test.shape[0],
y_test.shape[1]))

    # For each epoch...
    for epoch in tqdm(range(epochs)):
        # Shuffle the dataset
        idx = np.arange(X_train.shape[0])
        np.random.shuffle(idx)
        X_train_shuf = X_train[idx]
        y_train_shuf = y_train[idx]

        if opt == 'Adam':
            # 1. Forward pass
            y_pred_train = net.forward(X_train_shuf)
            # 2. Backward pass
            epoch_loss = net.backward(y_train_shuf)
            # 3. Update the weights
            net.update(lr=current_lr, b1=b1, b2=b2, eps=eps)

        elif opt == 'SGD': # mini-batch SGD
            num_train = X_train_shuf.shape[0]
            num_batches = (num_train + batch_size - 1) // batch_size
            epoch_loss = 0.0

            for b in range(num_batches):
                start = b * batch_size
                end = min(start + batch_size, num_train)
```

```

X_batch = X_train_shuf[start:end]
y_batch = y_train_shuf[start:end]

# 1. Forward pass
y_pred_train_batch = net.forward(X_batch)

# 2. Backward pass
train_loss_val = net.backward(y_batch)

# 3. Update the weights
net.update(lr=current_lr)

epoch_loss += train_loss_val

epoch_loss /= num_batches

train_loss[epoch] = epoch_loss

# Testing
y_pred_train = net.forward(X_train)
train_psnr[epoch] = psnr(y_train, y_pred_train)

y_pred_test = net.forward(X_test)
test_psnr[epoch] = psnr(y_test, y_pred_test)
predicted_images[epoch] = y_pred_test.copy()

return net, train_psnr, test_psnr, train_loss, predicted_images

```

Low Resolution Reconstruction

Low Resolution Reconstruction - SGD - None Mapping

```

# Display image
size = 32
train_data, test_data = get_image(size)

/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting
with ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
    img = imageio.imread(image_url)[..., :3] / 255.

```




```
B_dict = get_B_dict(size)

print('B_dict items:')
for k,v in B_dict.items():
    print('\t',k,np.array(v).shape)
```

```
B_dict items:
  none ()
  basic (2, 2)
  gauss_1.0 (16, 2)
```

```
# Set the hyperparameters
num_layers = 4
```

```

output_size = 3

# Adjust these
hidden_size = 256
epochs = 2000
learning_rate = 0.02



# Get input features
B_dict = get_B_dict(size=512)
X_train, y_train, X_test, y_test = get_input_features(B_dict,
mapping='none')
input_size = X_train.shape[1]

# Run NN experiment on input features
net, train_psnr, test_psnr, train_loss, predicted_images =
NN_experiment(
    X_train, y_train, X_test, y_test, input_size,
    num_layers,
    hidden_size, output_size, epochs, learning_rate,
    'mse', opt="SGD",
    batch_size=32)

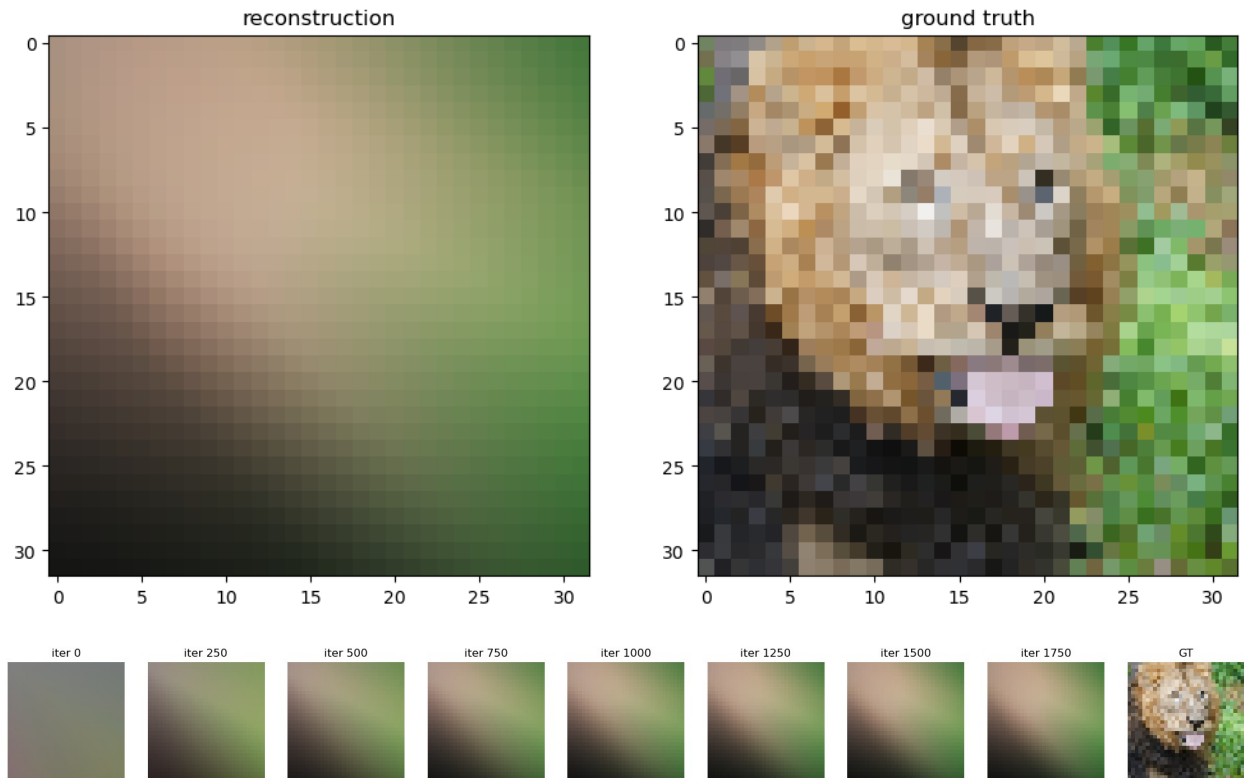
# Plot results of experiment
plot_training_curves(train_loss, train_psnr, test_psnr)
plot_reconstruction(net.forward(X_test), y_test)
plot_reconstruction_progress(predicted_images, y_test)

{"model_id": "9d0b159ccd9f496cbfebfff29eb34175", "version_major": 2, "version_minor": 0}

```



Final Test MSE 0.01968779947165195
Final Test psnr 14.047728270869522



Low Resolution Reconstruction - SGD - Various Input Mapping Strategies

```
def train_wrapper(mapping, size, num_layers, hidden_size, output_size,
epochs, learning_rate, loss_function, opt='SGD'):
```

```
    B_dict = get_B_dict(size)
    X_train, y_train, X_test, y_test =
get_input_features(B_dict,mapping)
    input_size = X_train.shape[1]

    net, train_psnrs, test_psnrs, train_loss, predicted_images =
NN_experiment(
                                X_train, y_train, X_test, y_test, input_size,
num_layers,
                                hidden_size, output_size, epochs, learning_rate,
loss_function, opt, batch_size=32)
```

```
    return {
        'net': net,
        'train_psnrs': train_psnrs,
        'test_psnrs': test_psnrs,
        'train_loss': train_loss,
        'pred_imgs': predicted_images
    }
```

```
outputs = {}
for k in tqdm(B_dict):
```

```

print("training", k)
outputs[k] = train_wrapper(k, size, num_layers, hidden_size,
output_size, epochs, learning_rate, 'mse', opt='SGD')

{"model_id": "099e8b8c4a3c425099b388ee9fad170b", "version_major": 2, "version_minor": 0}

training none

{"model_id": "3dc7c3e141e94231ad8931e52ed54478", "version_major": 2, "version_minor": 0}

training basic

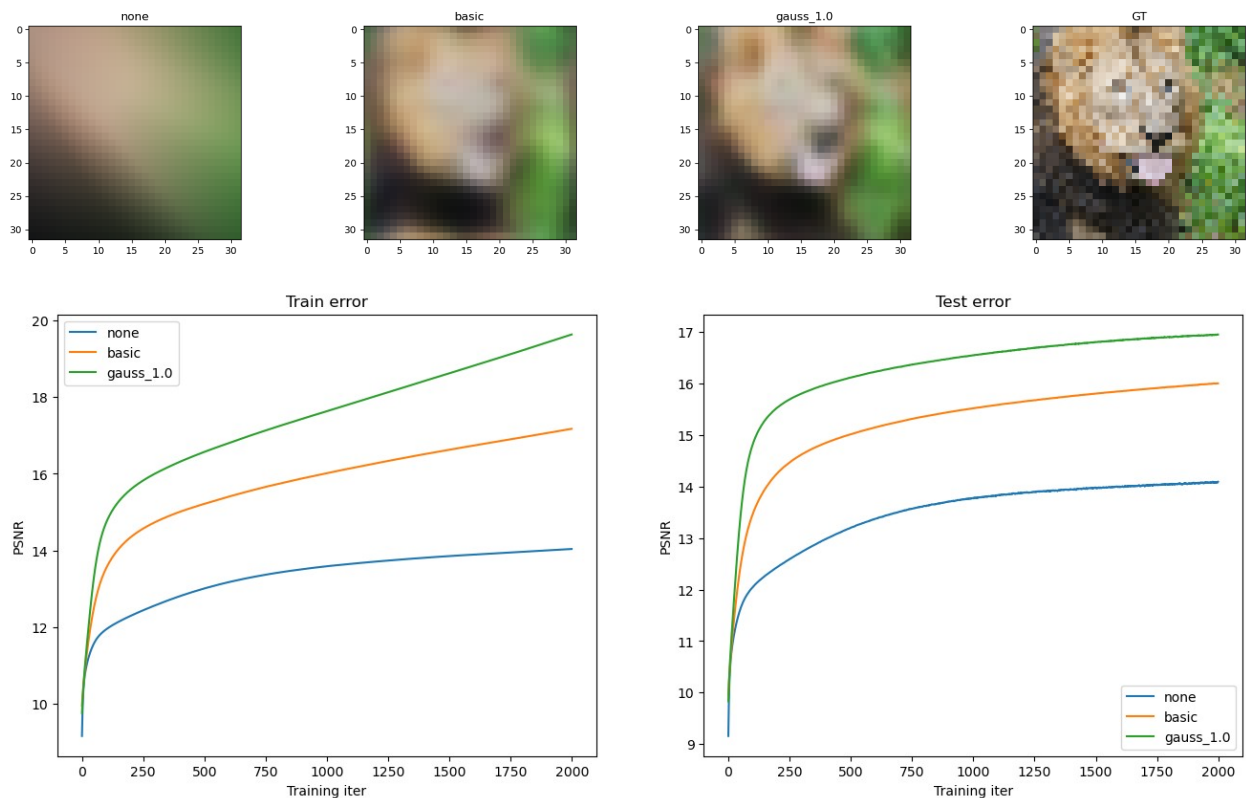
{"model_id": "2d4c1b64800d437f980e10aab16b5ca6", "version_major": 2, "version_minor": 0}

training gauss_1.0

{"model_id": "74467db8ee4e4d598632ca3b05b6a1dd", "version_major": 2, "version_minor": 0}

plot_feature_mapping_comparison(outputs, y_test.reshape(size,size,3))

```




High Resolution Reconstruction

High Resolution Reconstruction - SGD - Various Input Mapping Strategies

Repeat the previous experiment, but at the higher resolution.

```
# load hi-res image
size = 128
train_data, test_data = get_image(size)

/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting
with ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
  img = imageio.imread(image_url)[..., :3] / 255.
```



```
# Set the hyperparameters
num_layers = 4
output_size = 3

# Adjust these
hidden_size = 256
epochs = 300
learning_rate = 0.2

outputs = {}
for k in tqdm(B_dict):
    print("training", k)
```

```

    outputs[k] = train_wrapper(k,size, num_layers, hidden_size,
    output_size, epochs, learning_rate, 'mse', opt='SGD')

{"model_id":"7d15f574c1da4048bebd1144ad84bf59","version_major":2,"version_minor":0}

training none

{"model_id":"d276d18812c245648f309c327e9203ce","version_major":2,"version_minor":0}

training basic

{"model_id":"4771dd159f294b968f7e1e53209583fd","version_major":2,"version_minor":0}



training gauss_1.0

{"model_id":"36bdd029b5c84c6cb3b3225928aff0e2","version_major":2,"version_minor":0}

X_train, y_train, X_test, y_test =
get_input_features(get_B_dict(size), "none") # for getting y_test

plot_feature_mapping_comparison(outputs, y_test.reshape(size,size,3))

```




High Resolution Reconstruction - Image of your Choice

size = 128

```
train_data, test_data = get_image(size,  
image_url="https://i.pinimg.com/736x/0c/b4/6f/0cb46fc064b18cede9827908  
9d39e19b.jpg")
```

```
/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting  
with ImageIO v3 the behavior of this function will switch to that of  
iio.v3.imread. To keep the current behavior (and make this warning  
disappear) use `import imageio.v2 as imageio` or call  
`imageio.v2.imread` directly.
```

```
img = imageio.imread(image_url)[..., :3] / 255.
```



```

# Set the hyperparameters
num_layers = 4
output_size = 3

# Adjust these
hidden_size = 256
epochs = 300
learning_rate = 0.2

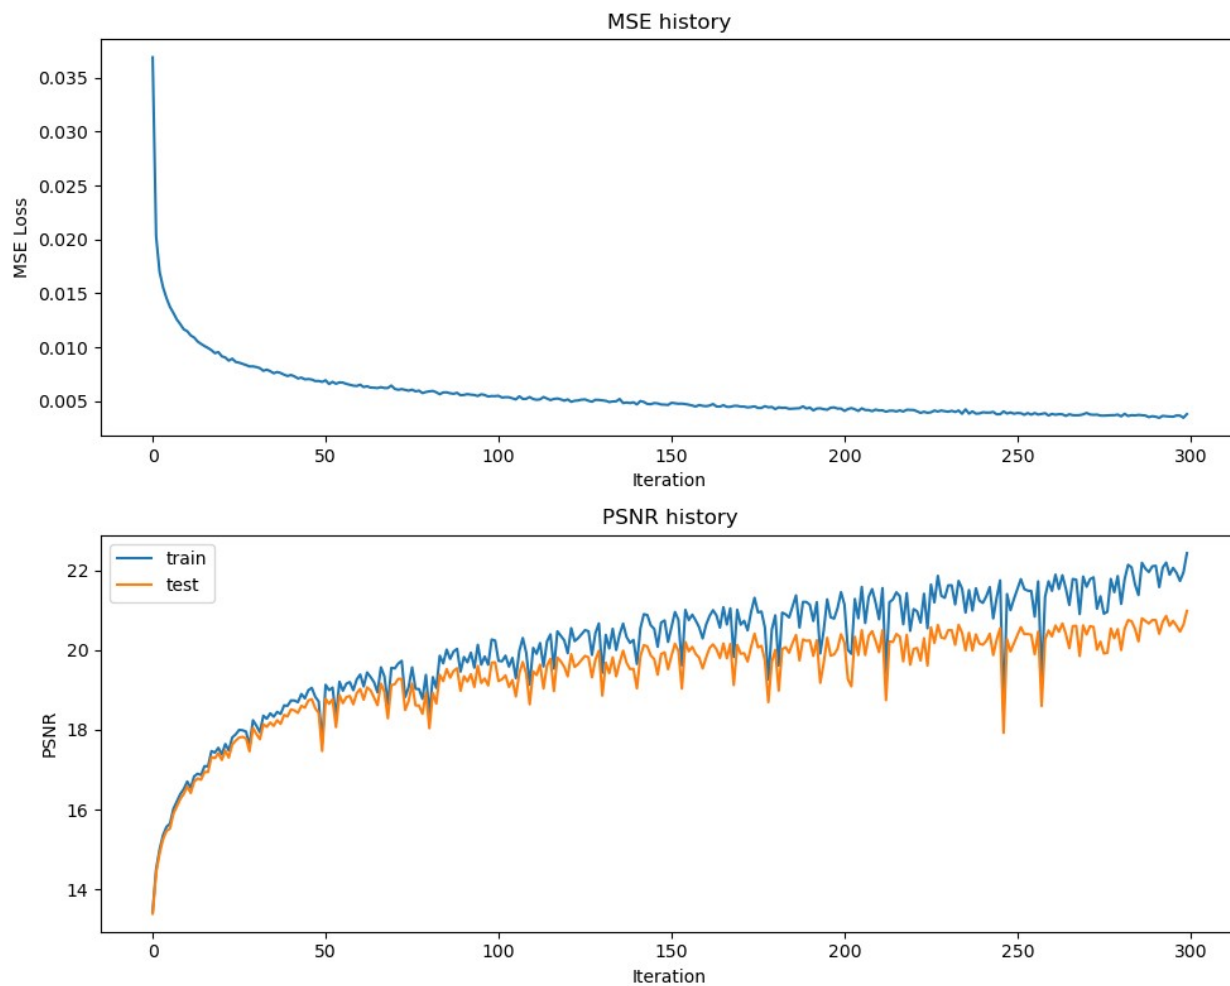
# Get input features
B_dict = get_B_dict(size)
X_train, y_train, X_test, y_test =
get_input_features(B_dict, mapping="gauss_1.0")
input_size = X_train.shape[1]

# Run NN experiment on input features
net, train_psnr, test_psnr, train_loss, predicted_images =
NN_experiment(
    X_train, y_train, X_test, y_test, input_size,
    num_layers,
    hidden_size, output_size, epochs, learning_rate,
    'mse', opt="SGD",
    batch_size=32)

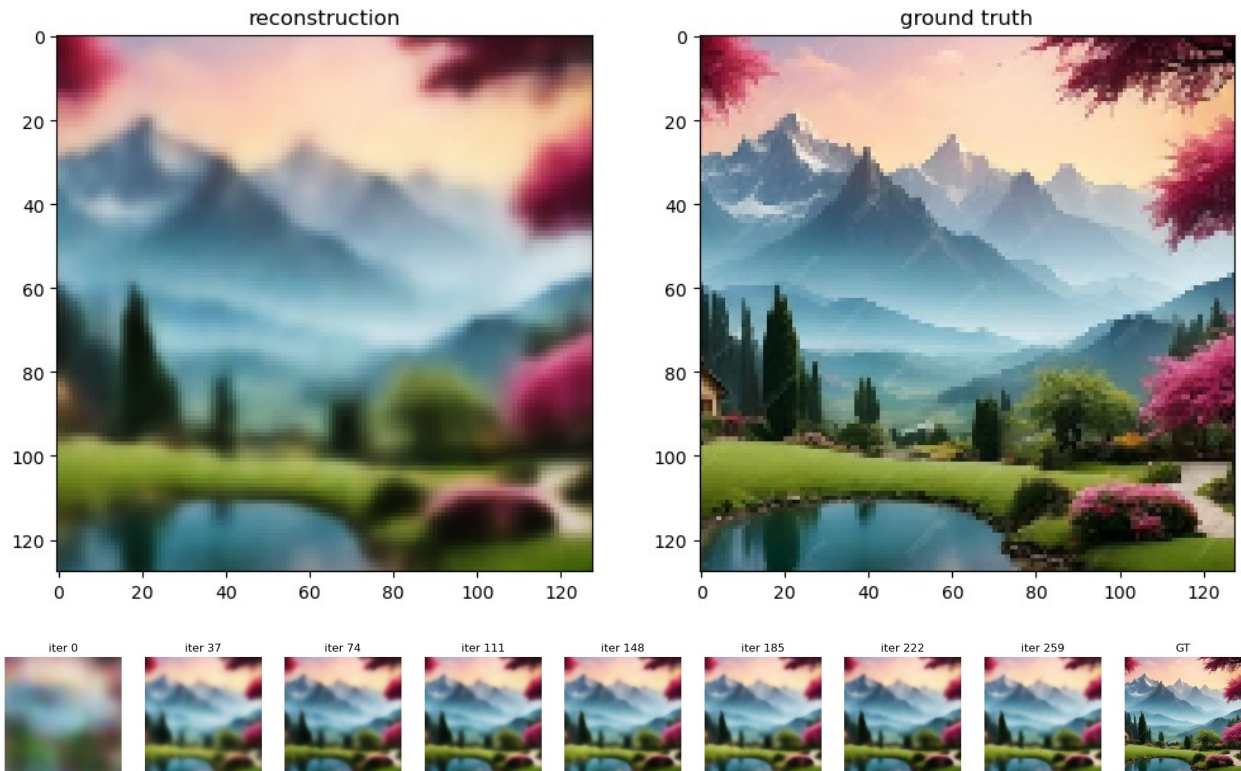
# Plot results of experiment
plot_training_curves(train_loss, train_psnr, test_psnr)
plot_reconstruction(net.forward(X_test), y_test)
plot_reconstruction_progress(predicted_images, y_test)

{"model_id": "f283deee84e54aed84524e013f6091eb", "version_major": 2, "version_minor": 0}

```



Final Test MSE 0.003990793849716292
Final Test psnr 20.97910710088206



1. Additional Experiments - Adam Optimizer

Low Resolution Reconstruction - Adam - None Mapping

```
# load low-res image
size = 32
train_data, test_data = get_image(size)

/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting
with ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
  img = imageio.imread(image_url)[..., :3] / 255.
```



```
# Set the hyperparameters
num_layers = 4
output_size = 3

# Adjust these
hidden_size = 256
epochs = 300
learning_rate = 0.0075

# get input features
B_dict = get_B_dict(size)
X_train, y_train, X_test, y_test =
get_input_features(B_dict, mapping='none')
```

```

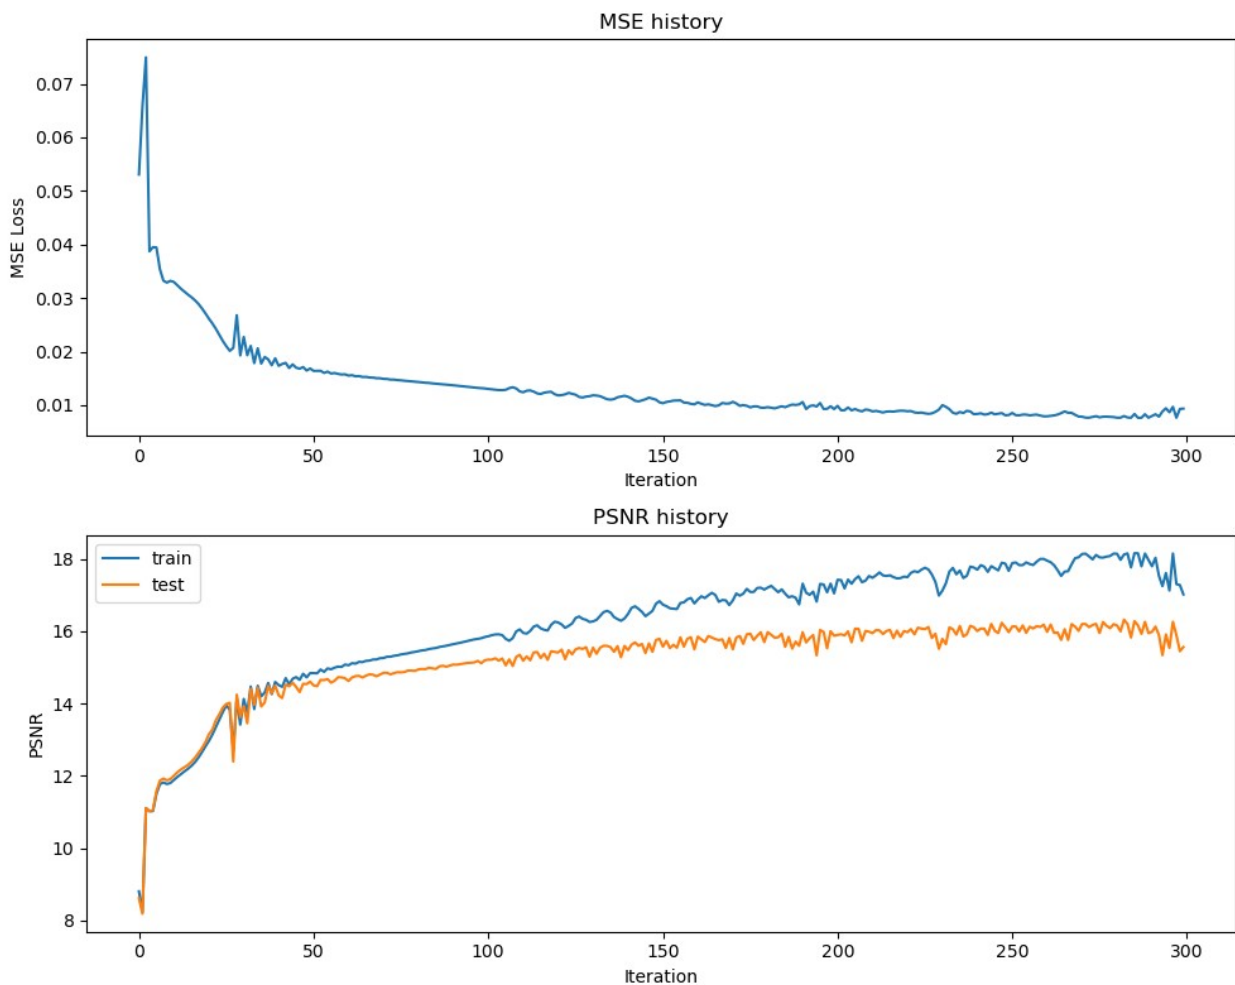
input_size = X_train.shape[1]

# run NN experiment on input features
net, train_psnr, test_psnr, train_loss, predicted_images =
NN_experiment(
    X_train, y_train, X_test, y_test, input_size,
    num_layers,
    hidden_size, output_size, epochs, learning_rate,
    'mse', opt="Adam",
    b1=0.9, b2=0.999, eps=1e-8)


# plot results of experiment
plot_training_curves(train_loss, train_psnr, test_psnr)
plot_reconstruction(net.forward(X_test), y_test)
plot_reconstruction_progress(predicted_images, y_test)

{"model_id": "94d167fabab440f1b98a85981fda2856", "version_major": 2, "version_minor": 0}

```



Final Test MSE 0.013894503187029448
Final Test psnr 15.56126981613346



Low Resolution Reconstruction - Adam - Various Input Mapping Strategies

```
# start training
outputs = {}
for k in tqdm(B_dict):
    print("training", k)
    outputs[k] = train_wrapper(k, size, num_layers, hidden_size,
    output_size, epochs, learning_rate, 'mae', opt='Adam')

{"model_id": "b7b21b9d07cf4eb7a604f95b95b1195c", "version_major": 2, "version_minor": 0}

training none

{"model_id": "bc31e1ece2264a8388acacfaf69264fe", "version_major": 2, "version_minor": 0}

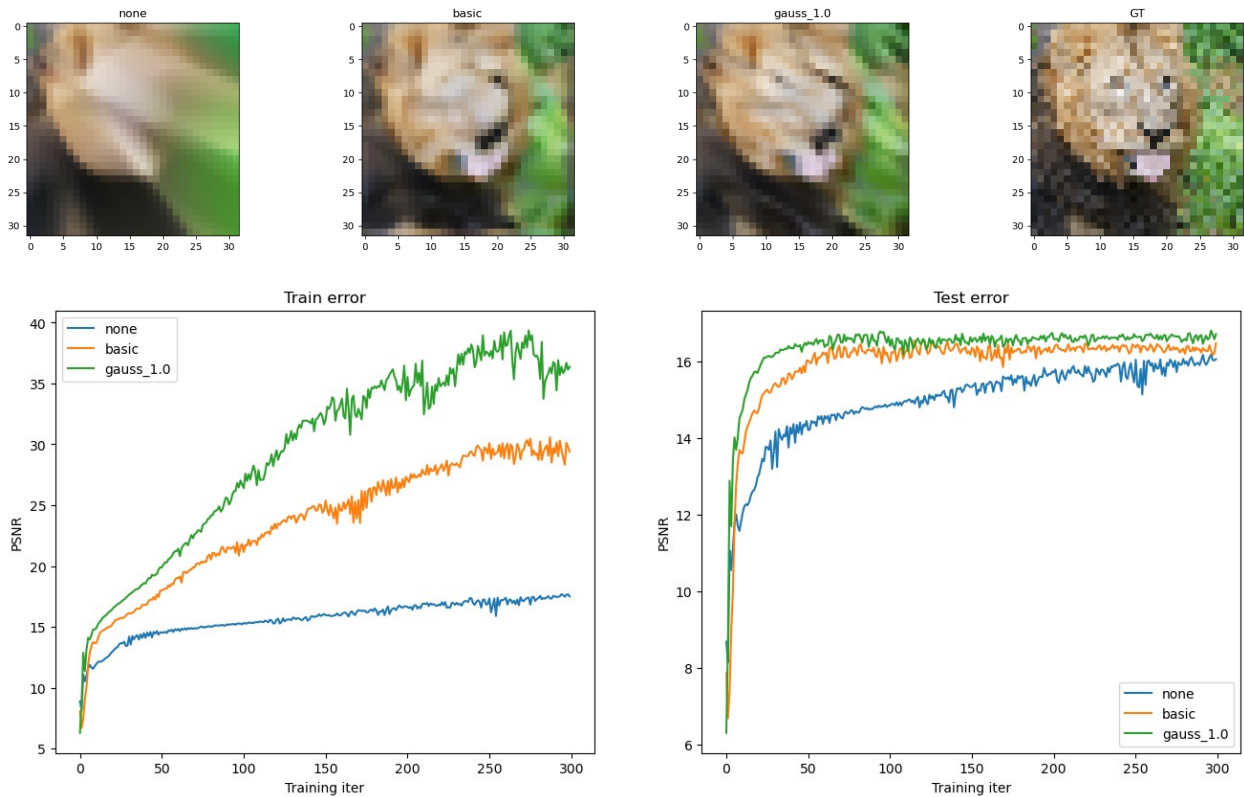
training basic

{"model_id": "6666d090ac084c4282a2be0243142746", "version_major": 2, "version_minor": 0}

training gauss_1.0
```

```
{"model_id": "c46641c9a4e546c18fd139f9c17dd9d2", "version_major": 2, "version_minor": 0}
```


```
plot_feature_mapping_comparison(outputs, y_test.reshape(size,size,3))
```



High Resolution Reconstruction - Adam - Various Input Mapping Strategies

```
# load hi-res image
size = 128
train_data, test_data = get_image(size)
```

```
/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting
with ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
img = imageio.imread(image_url)[..., :3] / 255.
```

```
# Set the hyperparameters
num_layers = 4
output_size = 3

# Adjust these
hidden_size = 256
epochs = 300
learning_rate = 0.0075

# start training
outputs = {}
for k in tqdm(B_dict):
    print("training", k)
```

```

    outputs[k] = train_wrapper(k, size, num_layers, hidden_size,
                                output_size, epochs, learning_rate, 'mse', opt='Adam')

{"model_id": "5e3e1ce35cdf4893ac36b3a9242e116f", "version_major": 2, "version_minor": 0}

training none

{"model_id": "c58274bc208a47eda3f440c1da833c82", "version_major": 2, "version_minor": 0}

training basic

{"model_id": "5f35dcd342f947a58c090d75a7f33bc6", "version_major": 2, "version_minor": 0}



training gauss_1.0

{"model_id": "0ef2cfad70a745a1ac8cbfe0cf8776f2", "version_major": 2, "version_minor": 0}

X_train, y_train, X_test, y_test =
get_input_features(get_B_dict(size), "none") # for getting y_test

plot_feature_mapping_comparison(outputs, y_test.reshape(size,size,3))

```



2. Additional Experiments - Deeper Network

Low Resolution Reconstruction

```
size = 32
train_data, test_data = get_image(size)

/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting
with ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
  img = imageio.imread(image_url)[..., :3] / 255.
```



```
num_layers = 6
output_size = 3

hidden_size = 512
epochs = 300
learning_rate = 0.2

outputs = {}
for k in tqdm(B_dict):
    print("training", k)
    outputs[k] = train_wrapper(k, size, num_layers, hidden_size,
                                output_size, epochs, learning_rate, 'mse', opt='SGD')
```

```
{"model_id": "669a9bb41a31461a89e44b5177761970", "version_major": 2, "version_minor": 0}
```

training none

```
{"model_id": "2fc977efee0e4c899add0fbe5c8725ef", "version_major": 2, "version_minor": 0}
```

training basic

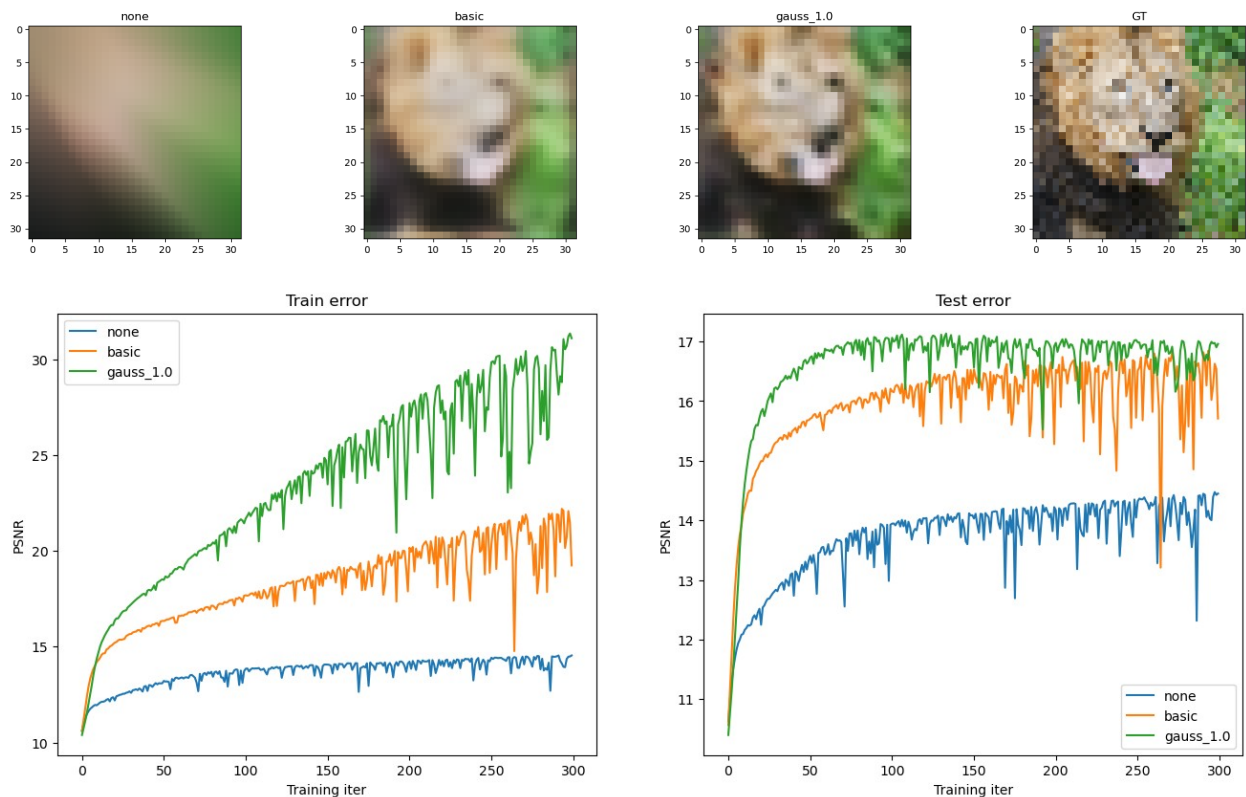
```
{"model_id": "3881f0c385c94f32ac9e026773d1df77", "version_major": 2, "version_minor": 0}
```

training gauss_1.0

```
{"model_id": "4a4b66bc5d094a16a54ffaaf51c83763", "version_major": 2, "version_minor": 0}
```

```
X_train, y_train, X_test, y_test =  
get_input_features(get_B_dict(size), "none") # for getting y_test
```


```
plot_feature_mapping_comparison(outputs, y_test.reshape(size,size,3))
```



High Resolution Reconstruction

```
size = 128
train_data, test_data = get_image(size)
```

```
/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting
with ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
  img = imageio.imread(image_url)[..., :3] / 255.
```



```
num_layers = 6
output_size = 3
```

```

# Adjust these
hidden_size = 768
epochs = 300
learning_rate = 0.2

outputs = {}
B_dict = get_B_dict(size)
for k in tqdm(B_dict):
    print("training", k)
    outputs[k] = train_wrapper(k, size, num_layers, hidden_size,
                                output_size, epochs, learning_rate, 'mse', opt='SGD')

{"model_id": "fa685c60588e4cf69e46f991d7053eea", "version_major": 2, "version_minor": 0}

training none

{"model_id": "6a1d5c46cc35482491acb38ef3c5ffcb", "version_major": 2, "version_minor": 0}

training basic





{"model_id": "f083cbb982454641bbdbbfa8c4b9c20c", "version_major": 2, "version_minor": 0}


training gauss_1.0

{"model_id": "b4f5d10a44ff4acdbff3b55b15284e71", "version_major": 2, "version_minor": 0}

X_train, y_train, X_test, y_test =
get_input_features(get_B_dict(size), "none")
plot_feature_mapping_comparison(outputs, y_test.reshape(size, size, 3))

```





3. Additional Experiments - Alternative Losses

```
size = 32
train_data, test_data = get_image(size)
num_layers = 3
output_size = 3
hidden_size = 256
epochs = 300
learning_rate = 0.2
```

```
outputs = {}
B_dict = get_B_dict(size=256)
k = 'gauss_1.0'
```

mae

```
outputs['mae'] = train_wrapper(k, size, num_layers, hidden_size,
output_size, epochs, learning_rate, 'mae', opt='SGD')
outputs['mse'] = train_wrapper(k, size, num_layers, hidden_size,
output_size, epochs, learning_rate, 'mse', opt='SGD')
outputs['huber'] = train_wrapper(k, size, num_layers, hidden_size,
output_size, epochs, learning_rate, 'huber', opt='SGD')
```

/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting with ImageIO v3 the behavior of this function will switch to that of `io.v3.imread`. To keep the current behavior (and make this warning disappear) use `import imageio.v2 as imageio` or call `imageio.v2.imread` directly.

```
img = imageio.imread(image_url)[..., :3] / 255.
```

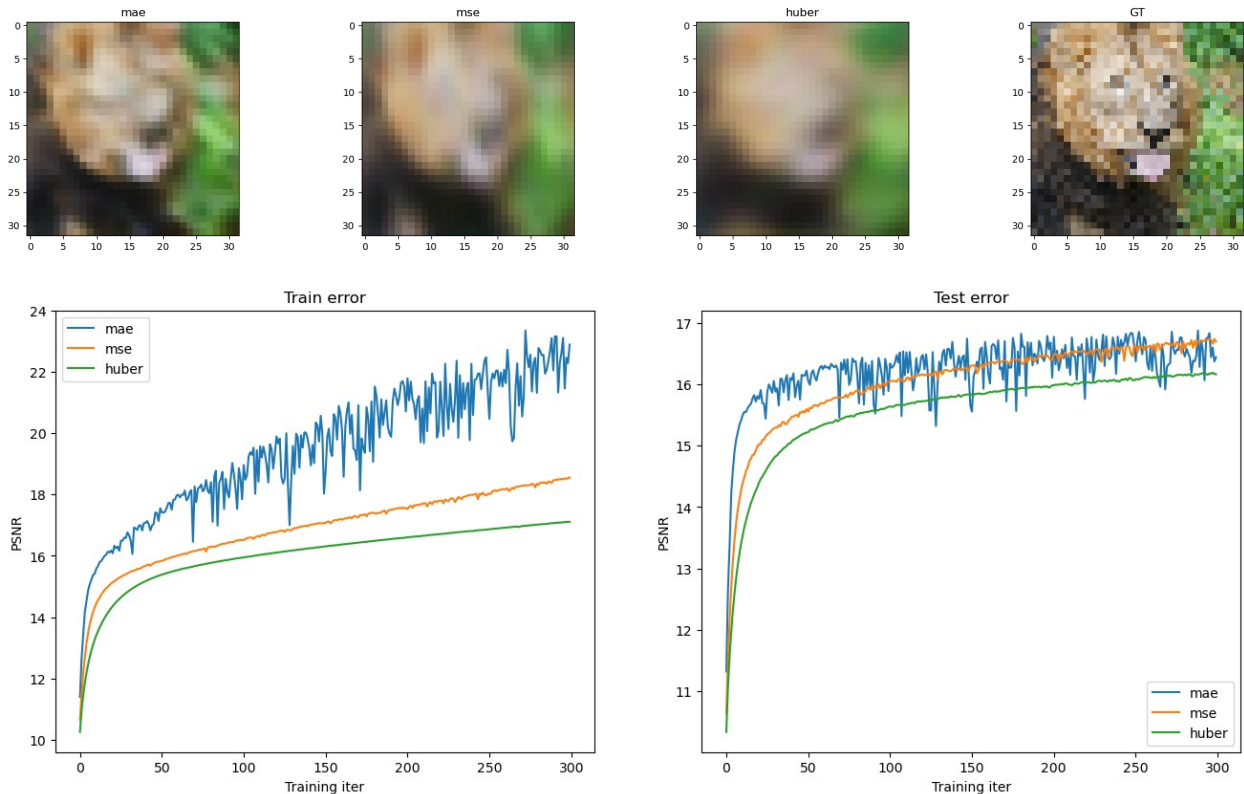



```
{"model_id":"cf1f5b8457e34f47a622977b4d194988","version_major":2,"version_minor":0}
```

```
{"model_id":"6ba55f8bf3bc49c6bbad9c2269885e6e","version_major":2,"version_minor":0}
```

```
{"model_id":"c0ef768e6c9f4d949d6038d1cfedde65","version_major":2,"version_minor":0}
```

```
X_train, y_train, X_test, y_test =  
get_input_features(get_B_dict(size), "none")  
plot_feature_mapping_comparison(outputs, y_test.reshape(size,size,3))
```



4. Additional Experiments - Gaussian Fourier Feature mapping hyperparameters

```
mapping_sizes = [8, 32, 128]
gaussian_scales = [1, 2]

size = 32

hyperparameter_configs = []
for scale in gaussian_scales:
    outputs = {}
    print('For mapping size:', scale)
    for m_size in mapping_sizes:
        B_dict = get_B_dict(scale, m_size)
        outputs[m_size] = train_wrapper('gauss_1.0', size, num_layers,
            hidden_size, output_size, epochs, learning_rate, 'mse', opt='SGD')




    X_train, y_train, X_test, y_test =
get_input_features(get_B_dict(size), "none")
    plot_feature_mapping_comparison(outputs,
y_test.reshape(size,size,3))

For mapping size: 1
```

```
{"model_id": "7d226ab076fd4dd4bcd94442c1cc8971", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "d6377b795dd54238860089d8f7be0c1d", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "49584400b2514af88cb4fccalc09374e", "version_major": 2, "version_minor": 0}
```

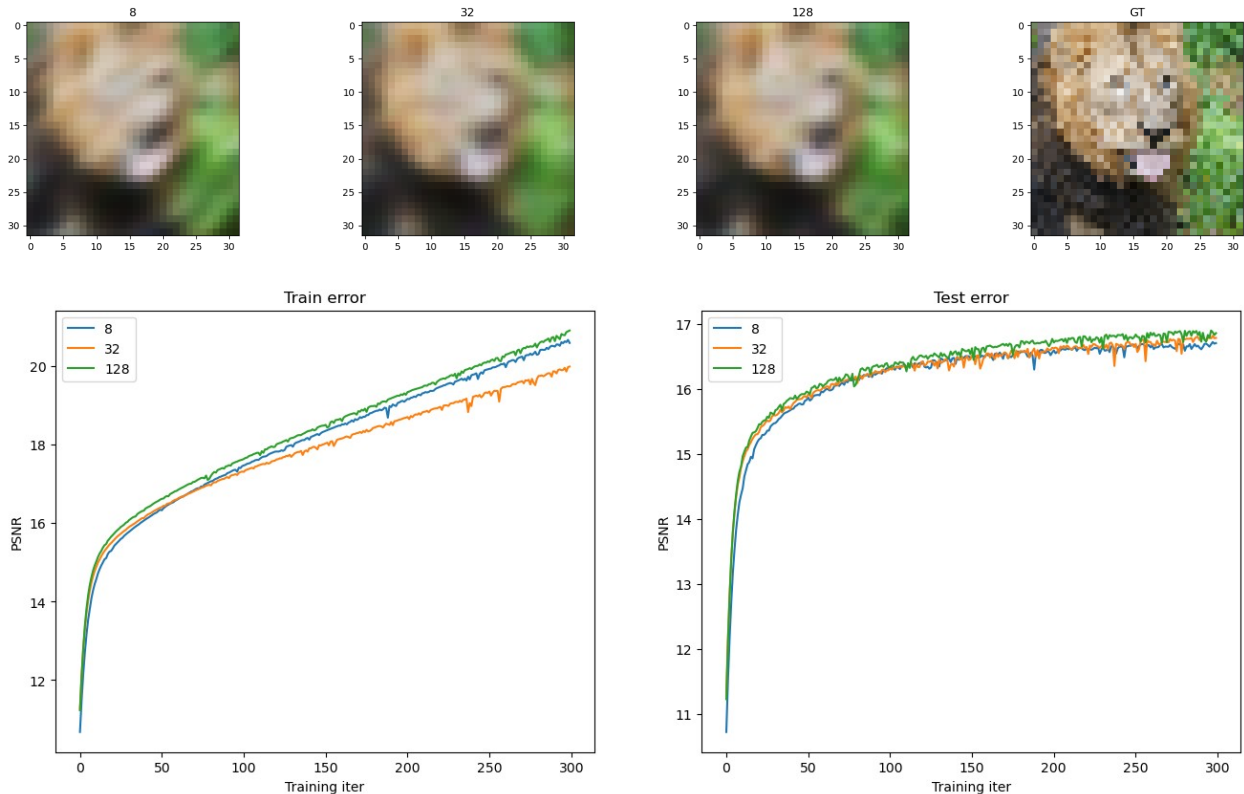


For mapping size: 2

```
{"model_id": "00265c56f4474408ad6d0774678ecd49", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "474f43abe59a43f9aa0fab088c784d50", "version_major": 2, "version_minor": 0}
```

```
{"model_id": "8454747e5c9e403489ab0cf6ab2819b4", "version_major": 2, "version_minor": 0}
```



5. Additional Experiments - Image inpainting or restoration scenarios

```
size = 128
def get_inpainting_data(size=512, mask_type='random', mask_ratio=0.1):

    train_data_orig, test_data = get_image(size)
    train_coords, train_img = train_data_orig
    test_coords, test_img = test_data

    train_size = train_img.shape[0] # Size of downsampled training
    image
    mask = np.ones((train_size, train_size), dtype=np.float32)

    if mask_type == 'random':
        random_mask = np.random.random((train_size, train_size)) >
mask_ratio
        mask = random_mask.astype(np.float32)

    elif mask_type == 'center':
        center = train_size // 2
        radius = int(train_size * np.sqrt(mask_ratio) / 2)
        y, x = np.ogrid[:train_size, :train_size]
```

```

        center_mask = (x - center)**2 + (y - center)**2 <= radius**2
        mask[center_mask] = 0.0

    elif mask_type == 'checkerboard':
        block_size = max(1, int(train_size * np.sqrt(mask_ratio) / 4))
        for i in range(0, train_size, block_size*2):
            for j in range(0, train_size, block_size*2):
                if i + block_size <= train_size and j + block_size <=
train_size:
                    mask[i:i+block_size, j:j+block_size] = 0.0

                if i + block_size*2 <= train_size and j + block_size*2
<= train_size:
                    mask[i+block_size:i+block_size*2,
j+block_size:j+block_size*2] = 0.0

    mask_3c = np.stack([mask] * 3, axis=-1)
    masked_img = train_img * mask_3c

    train_gt = train_img.copy()

    inpainting_train_data = [train_coords, masked_img, mask_3c]

    plt.figure(figsize=(15, 5))
    plt.subplot(1, 3, 1)
    plt.imshow(train_gt)
    plt.title('Original Image')

    plt.subplot(1, 3, 2)
    plt.imshow(mask, cmap='gray')
    plt.title('Mask')

    plt.subplot(1, 3, 3)
    plt.imshow(masked_img)
    plt.title('Masked Image')
    plt.show()

    return inpainting_train_data, test_data, train_gt

```

```

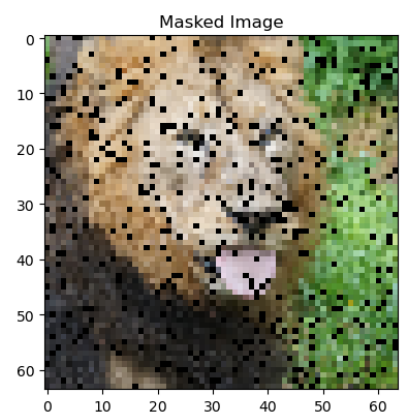



train_data, test_data, train_gt = get_inpainting_data(size)

```

```

/tmp/ipykernel_1671925/2979632171.py:6: DeprecationWarning: Starting
with ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
    img = imageio.imread(image_url)[..., :3] / 255.

```



```

num_layers = 4
output_size = 3

hidden_size = 256
epochs = 200
learning_rate = 0.2



B_dict = get_B_dict(size)
X_train, y_train, X_test, y_test =
get_input_features(B_dict, mapping="gauss_1.0")
input_size = X_train.shape[1]

net, train_psnr, test_psnr, train_loss, predicted_images =
NN_experiment(
    X_train, y_train, X_test, y_test, input_size,
    num_layers,
    hidden_size, output_size, epochs, learning_rate,
    'mse', opt="SGD",
    batch_size=32)

plot_training_curves(train_loss, train_psnr, test_psnr)
plot_reconstruction(net.forward(X_test), y_test)
plot_reconstruction_progress(predicted_images, y_test)

{"model_id": "0d854f96c4494bdaac17da1bf3160544", "version_major": 2, "version_minor": 0}

```



Final Test MSE 0.013294226003214523
Final Test psnr 15.753069467841119

